

Assignment 2

**Random Forests**

Sharin KG

14111033

## Question 1.

As discussed in class build a random forest classifier using bagging and random splitting using  $m$  random features to decide the split at each node. Find the number of trees at which the error levels off by using a binary search between 2 trees and 400-500 trees. Plot 5-fold cross-validated error rates against the number of trees in the forest and report the number at which error levels off.



## Solution

The number of trees at which the error levels off was found to be **93** by using Binary search between **2** and **500** trees. And the threshold for carrying out binary search was fixed to be **0.0015** based on the error rate at mid and last point.

```
#Function - Binary Search
binSearch <- function(h_error, low, high) {
  mid <- floor((low+high)/2)
  m_error <- getXError(mid)
  int = abs(h_error - m_error)
  if((abs(low-mid)<=1) || (abs(mid-high)<=1)) {
    return(high)
  }
  if(int<0.0015){
    high =mid
    return (binSearch(h_error, low, high))
  }
  else{
    low = mid
    return (binSearch(h_error, low, high))
  }
}

#Function - Validation Error
getXError <- function(noTree) {
  set.seed(2)
  error<-numeric(0)
  for(i in 1:5)
  {
    x_test_s = folds[[i]]
    x_train_s =setdiff(c(1:20000), x_test_s)
    x_test = Data[x_test_s,]
    x_train = Data[x_train_s,]
    actual = Data$lettr[x_test_s]
    x_fit <- randomForest(lettr ~ . ,data=x_train, ntree=noTree, replace = TRUE)
    predictions=predict(x_fit,newdata = x_test,type="class")
    error2=mean(actual!=predictions)
    error=rbind(error,error2)
  }
  return(mean(error))
}
```

Code Snippet 1. Function for Binary Search and Validation Error



**Question 2.**

Compute and report the out-of-bag error for the forest with the least error in part a).

**Solution**

The OOB for Number of Trees – 93 was found to be

```
set.seed(415)
rf=randomForest(lettr~.,data=Data,ntree = 93,replace = TRUE)
print(rf)
```

```
OOB estimate of error rate: 3.63%
```

*Code Snippet 2. Finding OOB Rate*

### Question 3.

Experiment with different values of  $m$  (say 1, 2, 4, 8) and report 5 fold cross-validated error rates in the form of a table fixing the number of trees at 1.25 times the number obtained in part a) (where the error levels off).

### Solution

Fixing the Number of Trees to be  $-93 \times 1.25 = 116$ , Cross Validated errors were generated.

As the number of attributes was 17, the least error rate was expected be around 4 ( $\sim \text{root}(17)$ ) and this was found to be true.

m Value	5 Fold Cross Validated Error Value (%)
1	5.60
2	3.87
4	3.72
8	4.16

Table 1.  $m$  vs. CV error rate

```
m=1
total_error<-array(0,dim=4)
for(k in 1:4){
  set.seed(415)
  cerror<-numeric(0)
  print("Count")
  print(m)
  for(i in 1:5) {
    cx_test_s = folds[[i]]
    cx_train_s = setdiff(c(1:20000), cx_test_s)
    cx_test = Data[cx_test_s,]
    cx_train = Data[cx_train_s,]
    cactual = Data$letter[cx_test_s]
    cx_fit <- randomForest(letter ~ . , data=cx_train, ntree=116, replace = TRUE, mtry = m)
    prediction=predict(cx_fit,newdata = cx_test,type="class")
    cerror=rbind(cerror,mean(cactual!=prediction))
  }
  print(cerror)
  total_error[k]=mean(cerror)
  m=m*2
}
print(total_error*100)
```

Code Snippet 3. Generating CV Error rates for  $m=1,2,3,4$

### Question 4.

Study the effect of the size of the randomly sampled data set from  $L$  while constructing a tree. Start by sampling 10% of the points from  $L$  for constructing a tree and go up to 80% in increments of 10% and the number of trees, say  $T$ , in the forest at which the error levels off in each case. Find the 5 fold cross-validated error for a forest with  $T$  trees and plot it against the size of the sampled data set expressed as a percentage of  $L$ . Comment on the results and say whether bagging is justified as a randomization method to select samples?

### Solution

For sampling size ranging from 10% to 80% of the training set, the CV error rate vs. Size of sampled data plots were made, after finding the error off number of trees for each time.

- **CV Error vs. No of Trees Plot**

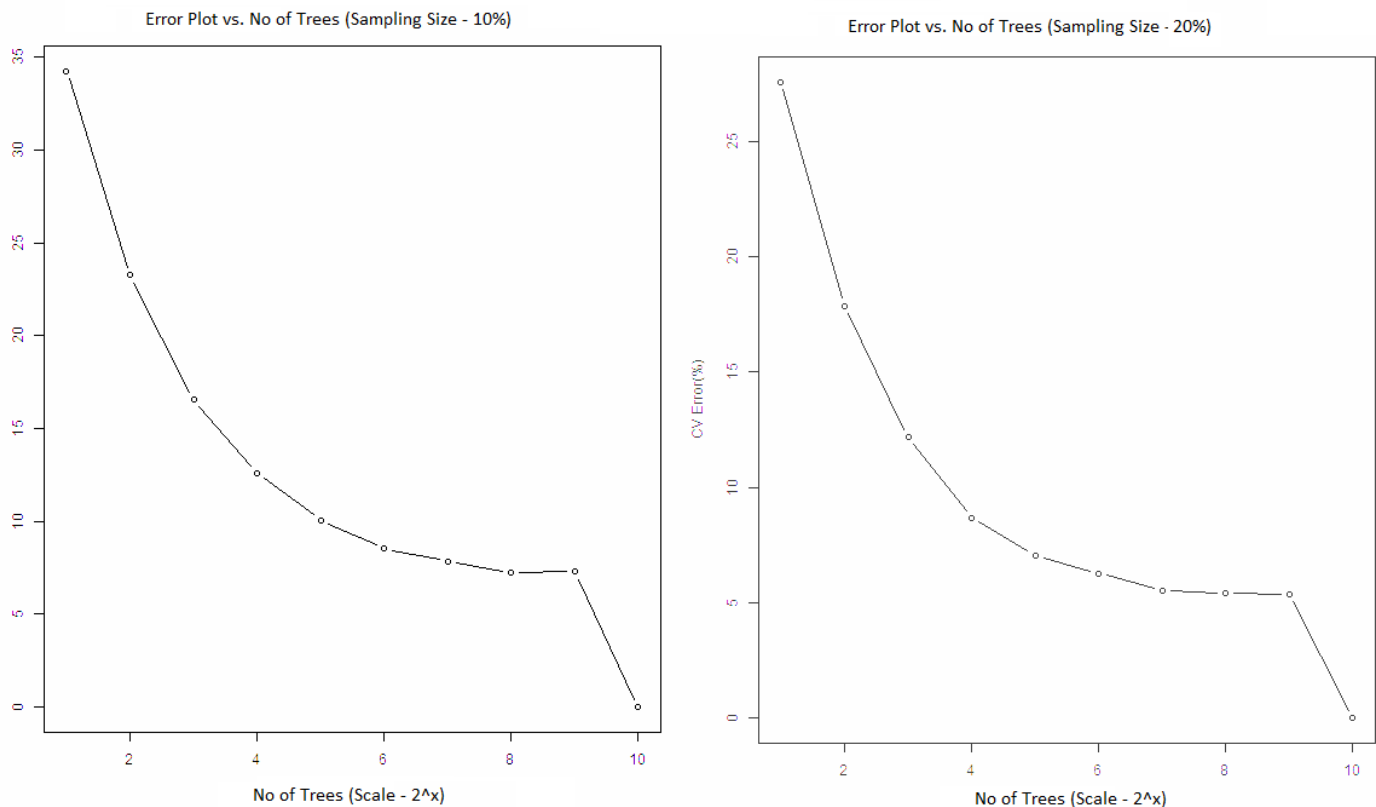


Figure 2: For Sampling Rates 10 and 20%

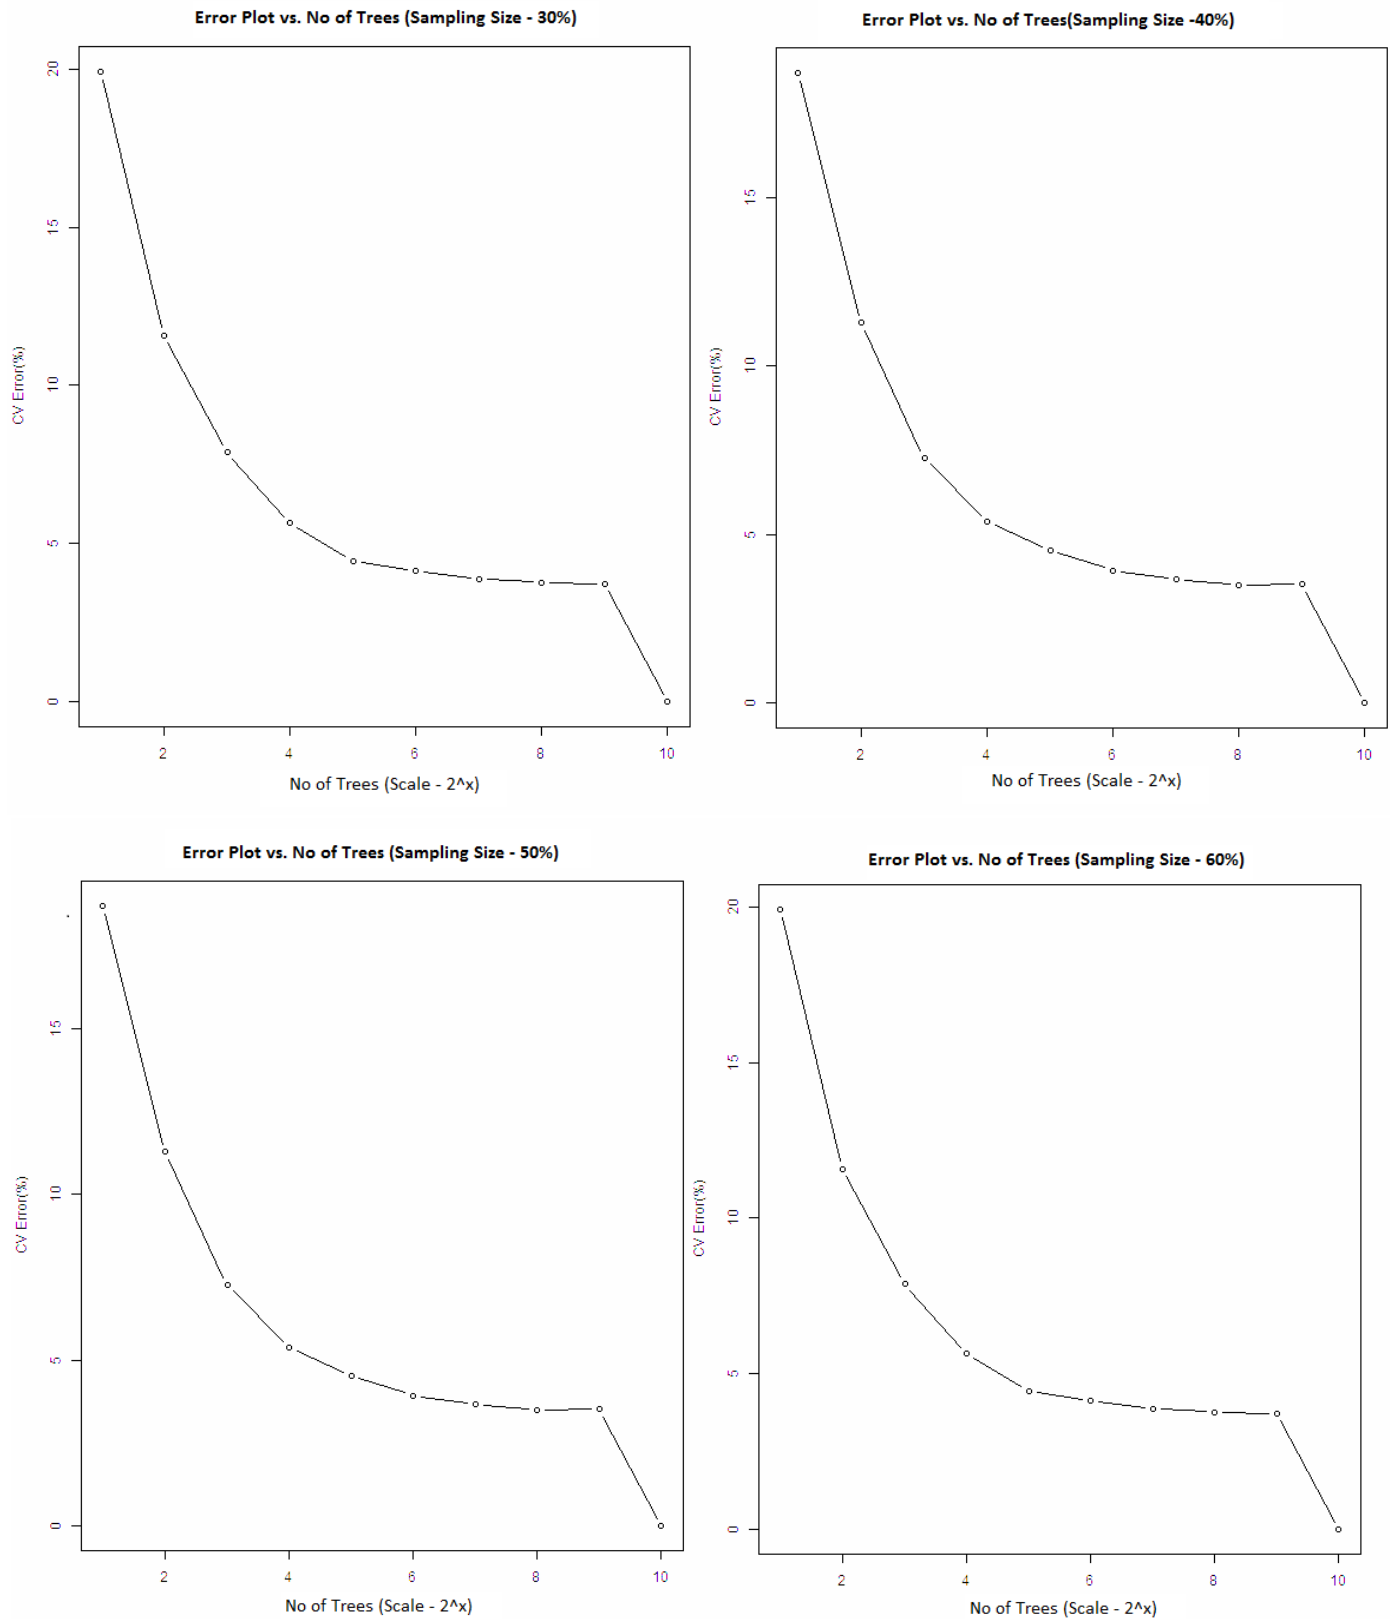


Figure 2: For Sampling Rates 30 -60%





Figure 4: For Sampling Rates 70 and 80%

From the above and below graphs, we can clearly see the trend that is – As the size of the sampled data from the Learning Set increases, the error rate clearly decreases up to a certain point after which it becomes stable (around 60%), and since bagging also takes into account ~63% of randomized data, it is justified.



- 5 fold CV vs. Size of the Sampled Data from  $L$



```

#Function - Validation Error
getXError <- function(noTree,size){
  set.seed(415)
  error<-numeric(0)
  for(i in 1:5){
    x_test_s = folds[[i]]
    x_train_s =setdiff(c(1:20000), x_test_s)
    x_test = Data[x_test_s,]
    x_train = Data[x_train_s,]
    actual = Data$lettr[x_test_s]
    x_fit <- randomForest(lettr ~ . ,data=x_train, ntree=noTree,sampsize =
nrow(x_train)*size,replace = FALSE)
    predictions=predict(x_fit,newdata = x_test,type="class")
    error=rbind(error,mean(actual!=predictions))
  }
  return(mean(error))
}

#Function - Binary Search
binSearch <- function(h_error,low,high,ssize){
  mid <- floor((low+high)/2)
  m_error <- getXError(mid,ssize)
  int = abs(h_error - m_error)
  if((abs(low-mid)<=1) || (abs(mid-high)<=1)){
    return(high)
  }
  if(int<0.0015){
    high =mid
    return (binSearch(h_error,low,high,ssize))
  }
  else{
    low = mid
    return (binSearch(h_error,low,high,ssize))
  }
}

arr_opt_pt<-array(0,dim=8)
arr_x<-array(0,dim=8)
for(k in 1:8){
  s_size = sample_size[k]
  tresh = getXError(500,s_size)
  arr_opt_pt[k]=binSearch(tresh,2,500,s_size)
  arr_x[k]<- getXError(arr_opt_pt[k],s_size)
  error_plot=array(0,dim=10)
  for(i in 1:9){
    if(i<9){
      error_plot[i]=getXError(2^i,s_size)
    }
    else{
      error_plot[i]=getXError(500,s_size)
    }
  }
  fileName = paste("Plot ", k, ".jpg")
  png(filename = fileName, width = 600, height = 700)
  plot(error_plot*100,type = "b",xlab = "No. of Trees (Scale - 2^x)",ylab = "CV Error(%)", main
= "Error Plot vs No. of Trees")
  dev.off()
}
png(filename = "Part4.jpg", width = 600, height = 700)
plot(x=sample_size, y=arr_x*100,type = "b",xlab = "Sampling Size (%)",ylab = "CV Error(%)",
main="Error Off vs Sample Size")
dev.off()

```

Code Snippet for Part d of the Question